

Shree Prashant Joshi

B.Tech Computer Science · 3rd Year · Pune, India

☎ +91 8446100247 | ✉ Shreejoshi1805@gmail.com | [LinkedIn](#) | [GitHub](#)

PROFESSIONAL SUMMARY

Driven 3rd-year Computer Science student with practical experience in Full-Stack Web Development, Machine Learning, and Python-based Data Analytics. Hands-on builder with 4+ projects spanning real-time applications, ML pipelines, and AI-powered systems. Currently exploring Large Language Models (LLMs) and Retrieval-Augmented Generation (RAG). Seeking a challenging internship to contribute and grow.

EDUCATION

Pimpri Chinchwad University (PCU) – School of Engineering & Technology 2023 – 2027 (Expected)

B.Tech – Computer Engineering | CGPA: 7.47 / 10

Abhishek Vidyalayam

Higher Secondary (Class 12 – Science) | 2023

Abhishek Vidyalayam School, Pune

10th Standard | MSBSHSE | 2021

TECHNICAL SKILLS

Languages:	C, Python, JavaScript, HTML, CSS
Python / Data:	Pandas, NumPy, Matplotlib, Scikit-learn, Jupyter Notebook
ML / AI:	Supervised & Unsupervised Learning, Classification, Regression, Decision Tree, Random Forest, Model Evaluation (Accuracy, Precision, Recall), RAG, LLM, Vector Embedding
Databases:	MongoDB, MySQL, Vector DB
CS Fundamentals:	Operating Systems, Computer Networks, DBMS, Object-Oriented Programming (OOP)
Tools:	Git, GitHub, VS Code, Postman, Eclipse, IntelliJ IDEA, REST API, Docker

PROJECTS

RAG PDF Assistant · [Python](#) · [LangChain](#) · [FAISS](#) · [LLM](#) · [Vercel](#)

- Built a RAG pipeline that extracts, chunks, and embeds document text into a vector database (FAISS/Chroma) for semantic similarity search.
- Integrated Ollama (Llama 3.2:3B) locally for LLM-based answer generation using retrieved context chunks via a Flask REST API, reducing hallucination and improving response accuracy.
- Built an interactive web UI (HTML/JS) with Flask backend, enabling real-time Q&A over uploaded documents through /ask API endpoint.
- Deployed the full application to Vercel with CI/CD via GitHub Actions workflow, making it publicly accessible.

Tech Stack: Python, LangChain, FAISS, Ollama/Embeddings, Flask/FastAPI, React.js, Vercel

Multi Agent Research System · [Python](#) · [LangChain](#) · [LangGraph](#) · [Gemini](#) · [Docker](#)

- Built an autonomous multi-agent AI research pipeline using LangChain and LangGraph with four specialized agents — Search, Reader, Writer, and Critic — connected via stateful graph-based orchestration.
- Search Agent retrieves relevant URLs; Reader Agent scrapes and extracts clean content using BeautifulSoup; Writer Agent generates structured reports via Google Gemini API; Critic Agent reviews and refines the final output.
- Implemented modular agent architecture with each agent independently defined, enabling easy extension and support for additional LLM providers (OpenAI, Claude, Mistral).
- Containerized the full pipeline using Docker for consistent, reproducible deployments across environments.

Tech Stack: Python, LangChain, LangGraph, Google Gemini API, BeautifulSoup4, Requests, Docker, python-dotenv

CERTIFICATIONS

- Red Hat Linux Fundamentals – Linux Commands, File System, Permissions & Shell Basics
- IBM Machine Learning – Supervised and Unsupervised Learning